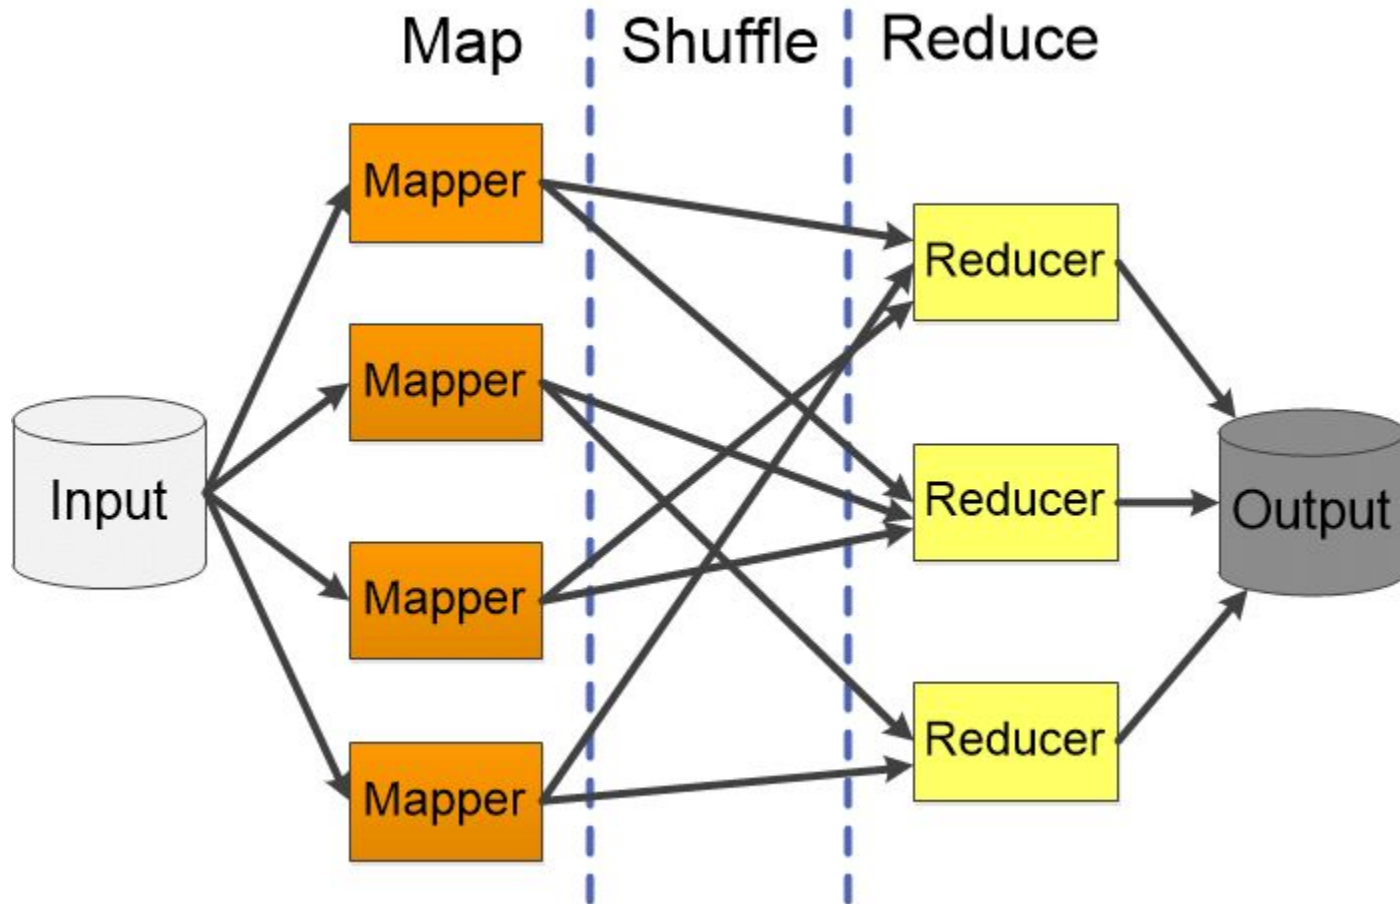


Introduction

- MapReduce is a **software framework for processing (large) data sets in a distributed fashion over a several machines.**
- The core idea behind MapReduce is mapping your data set into a collection of <key, value> pairs, and then reducing over all pairs with the same key.
- MapReduce is a **programming paradigm that enables massive scalability across hundreds or thousands of servers in a Hadoop cluster.**
- As the processing component, MapReduce is the heart of Apache Hadoop.
- The term "MapReduce" refers to two separate and distinct tasks that Hadoop programs perform.
- MapReduce program executes in three stages, namely **map stage, shuffle stage, and reduce stage.**

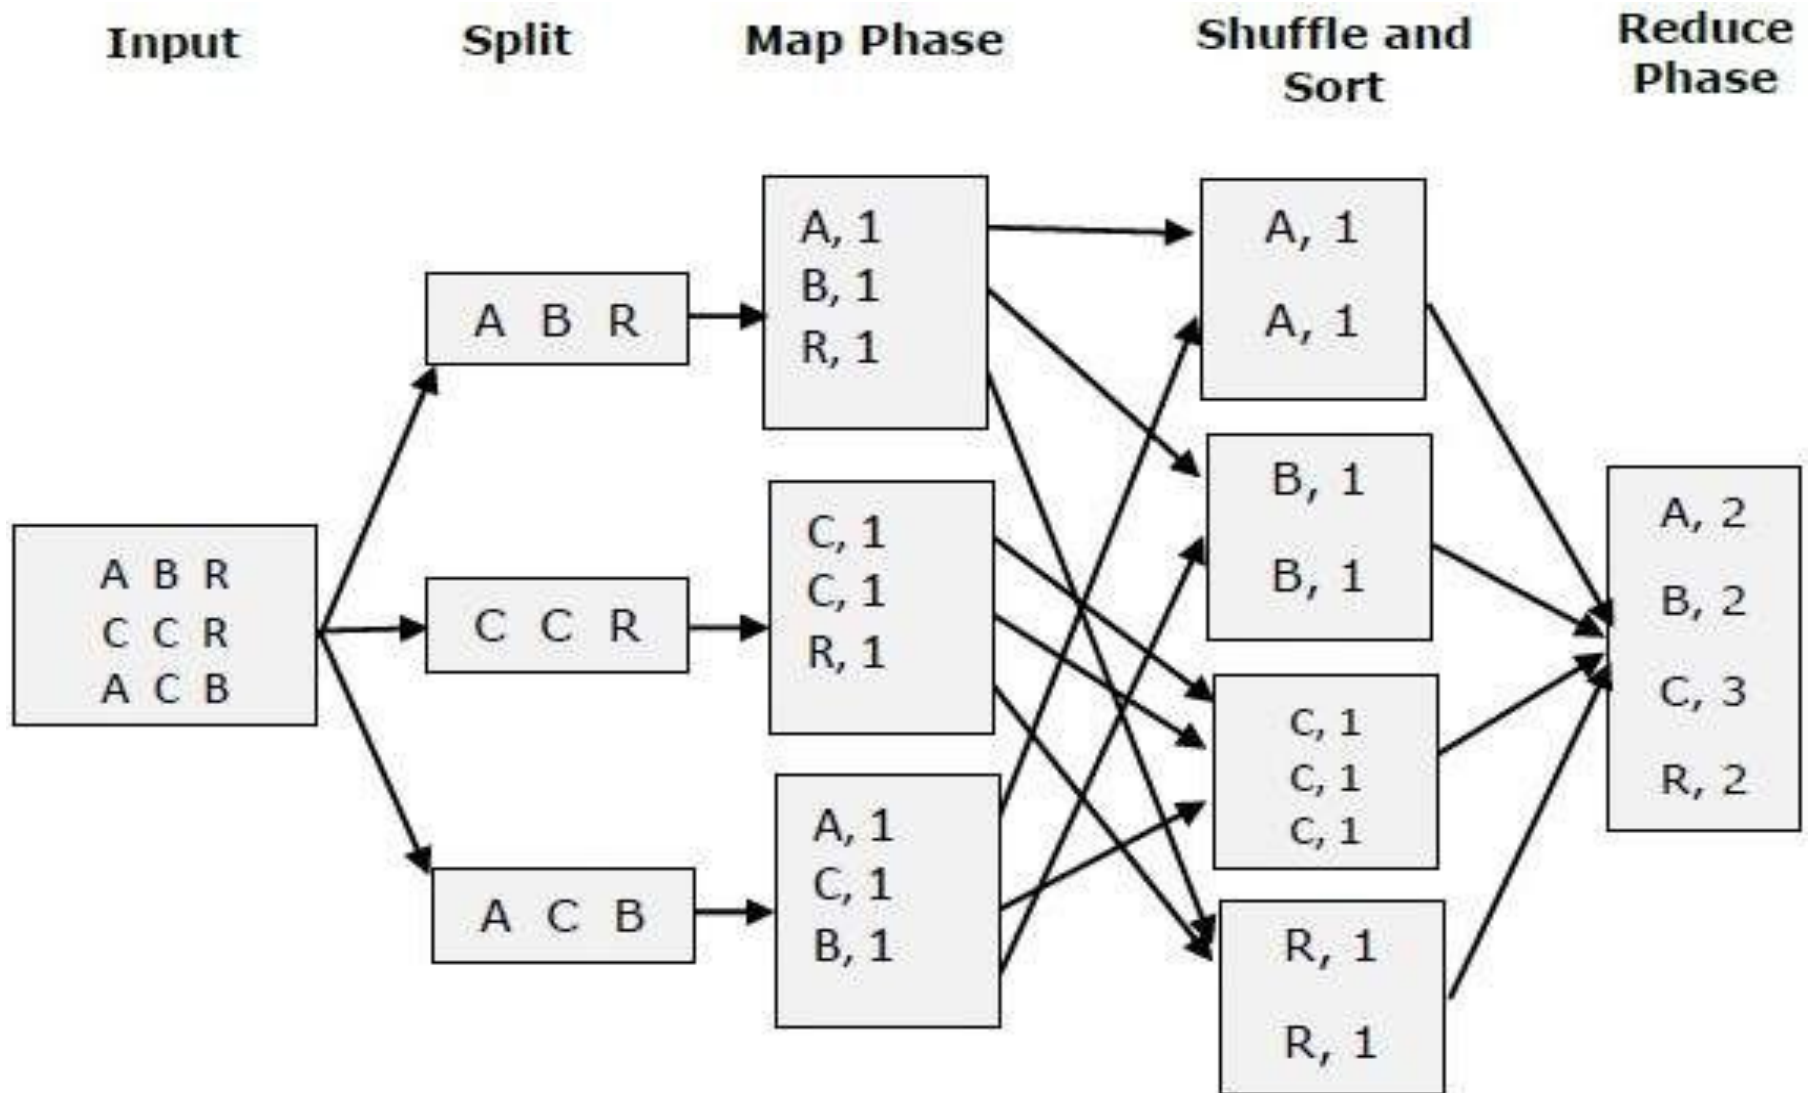
MapReduce Diagram



Task for MapReduce

- The MapReduce algorithm contains two important tasks, namely **Map and Reduce**.
- The Map task takes a set of data and converts it into another set of data, where individual elements are broken down into tuples(key-value pairs).

MapReduce - Example



Input phase

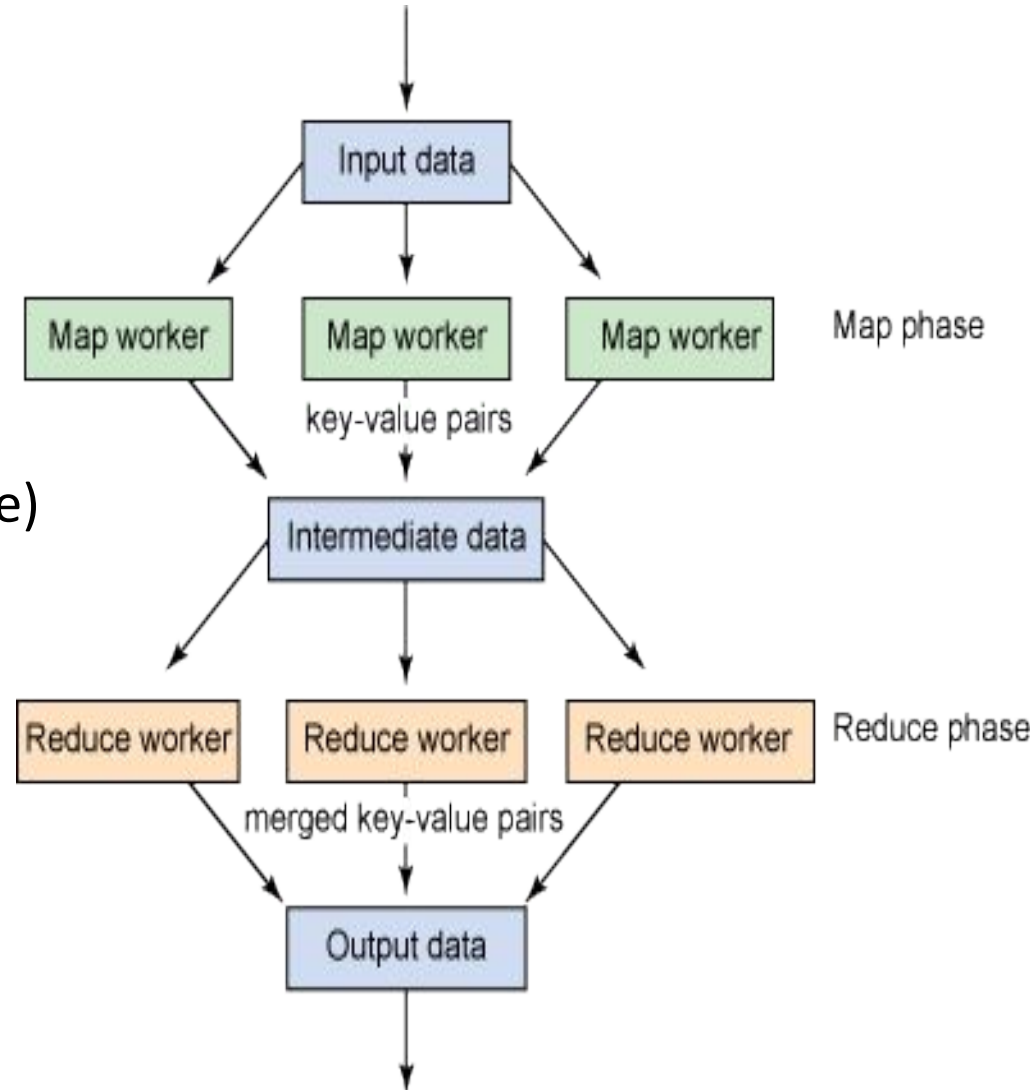
- Map task takes care of loading, parsing, transforming and filtering.
- The responsibility of reduce task is grouping and aggregating data that is produced by map tasks to generate final output.
- Each map task is broken into the following phases :
 - 1. RecordReader
 - 2. Mapper
 - 3. Combiner
 - 4. Partitioner

Output phase

- The output produced by map task is known as intermediate keys and values.
- These intermediate keys and values are sent to reducer.
- The reducer tasks are broken into the following phases :
 - 1. Shuffle
 - 2. Sort
 - 3. Reducer
 - 4. output Format

Major Components

- User Components:
 - Mapper
 - Reducer
 - Combiner (Optional)
 - Partitioner (Optional) (Shuffle)
 - Writable(s) (Optional)
- System Components:
 - Master
 - Input Splitter
 - Output Committer



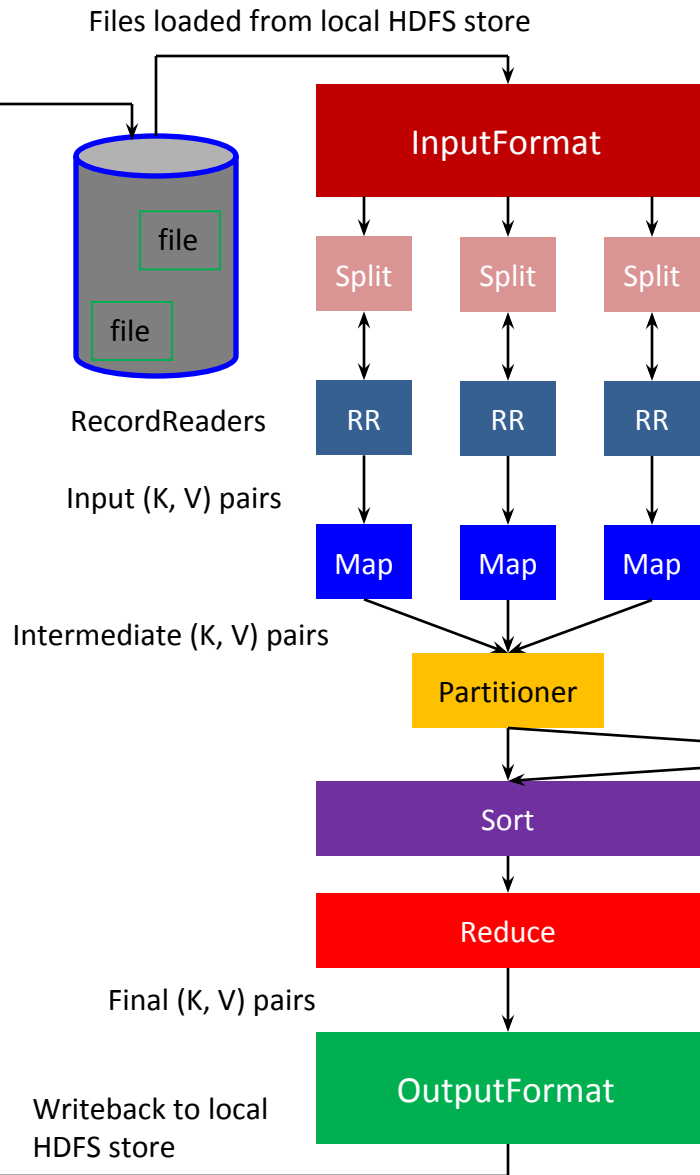
Mapper

- A mapper maps the input key-value pairs into a set of intermediate key-value pairs.
- Maps are individual tasks that have the responsibility of transforming input records into intermediate key-value pairs.
- Each map task is broken into the following phases :
 - 1. RecordReader
 - 2. Mapper
 - 3. Combiner
 - 4. Partitioner

Hadoop MapReduce: A Closer Look

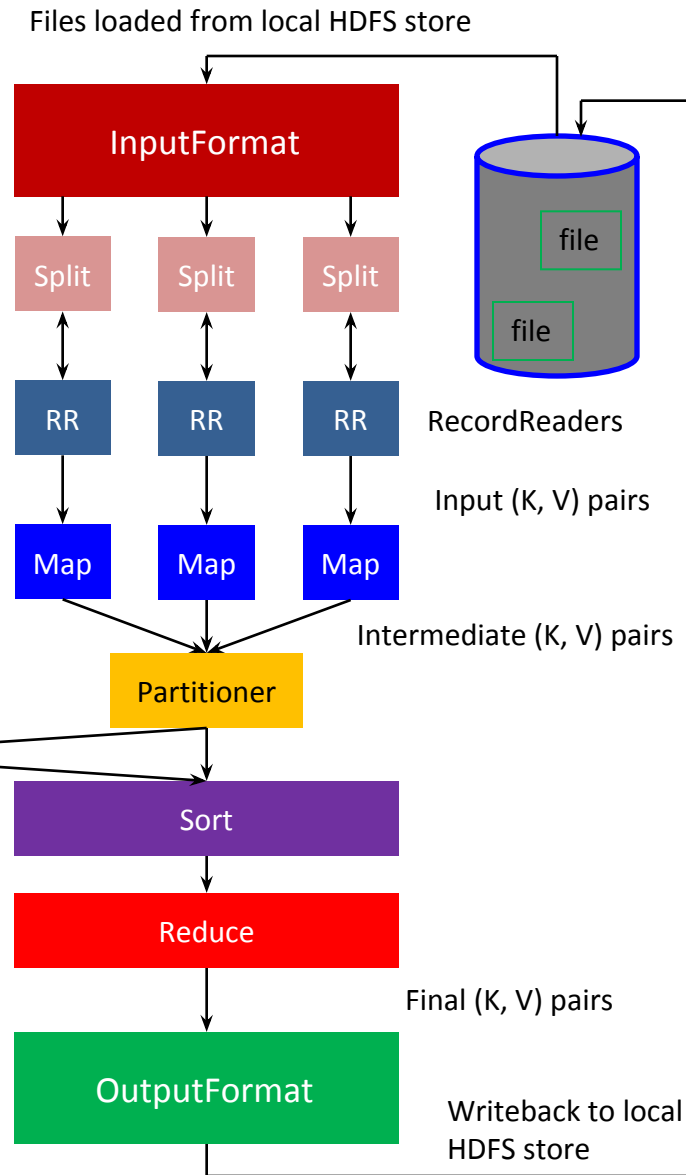
Node 1

Node 2




Shuffling Process

Intermediate (K,V) pairs exchanged by all nodes




Input Files

- *Input files* are where the data for a MapReduce task is initially stored
- The input files typically reside in a distributed file system (e.g. HDFS)
- The format of input files is arbitrary
 - Line-based log files
 - Binary files
 - Multi-line input records
 - Or something else entirely




InputFormat

- How the input files are split up and read is defined by the *InputFormat*
- InputFormat is a class that does the following:
 - Selects the files that should be used for input
 - Defines the *InputSplits* that break a file
 - Provides a factory for *RecordReader* objects that read the file



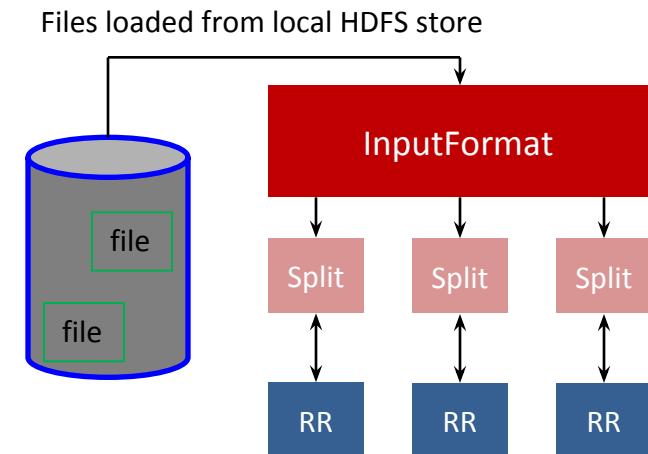
Input Splits

- An *input split* describes a unit of work that comprises a single map task in a MapReduce program
- By default, the InputFormat breaks a file up into 64MB splits
- By dividing the file into splits, we allow several map tasks to operate on a single file in parallel
- If the file is very large, this can improve performance significantly through parallelism
- Each map task corresponds to a *single* input split



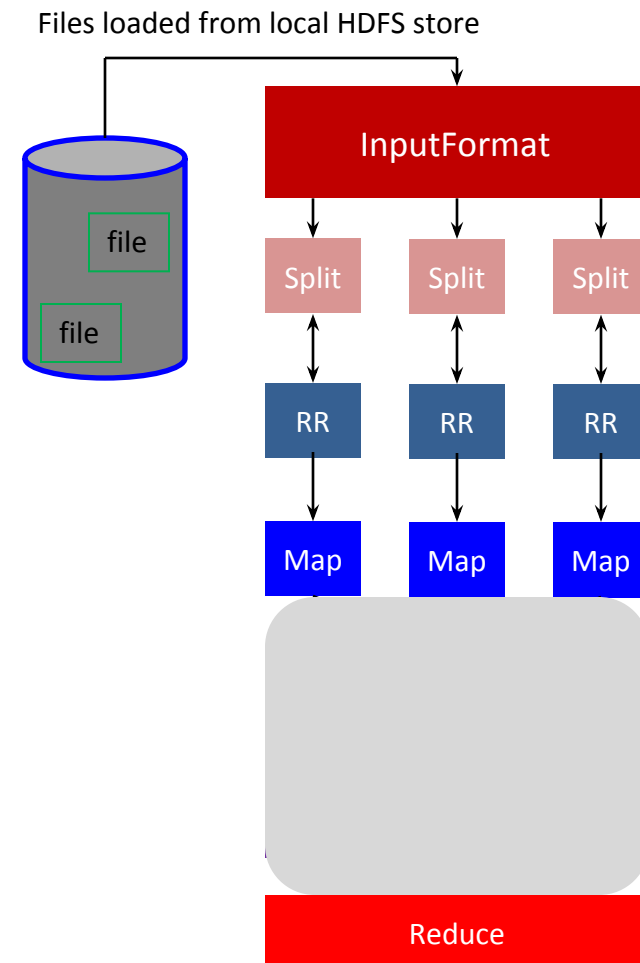
RecordReader

- The input split defines a slice of work but does not describe how to access it
- The *RecordReader* class actually loads data from its source and converts it into (K, V) pairs suitable for reading by Mappers
- The RecordReader is invoked repeatedly on the input until the entire split is consumed
- Each invocation of the RecordReader leads to another call of the map function defined by the programmer



Mapper and Reducer

- The *Mapper* performs the user-defined work of the first phase of the MapReduce program
- A new instance of Mapper is created for each split
- The *Reducer* performs the user-defined work of the second phase of the MapReduce program
- A new instance of Reducer is created for each partition
- *For each key in the partition assigned to a Reducer, the Reducer is called once*



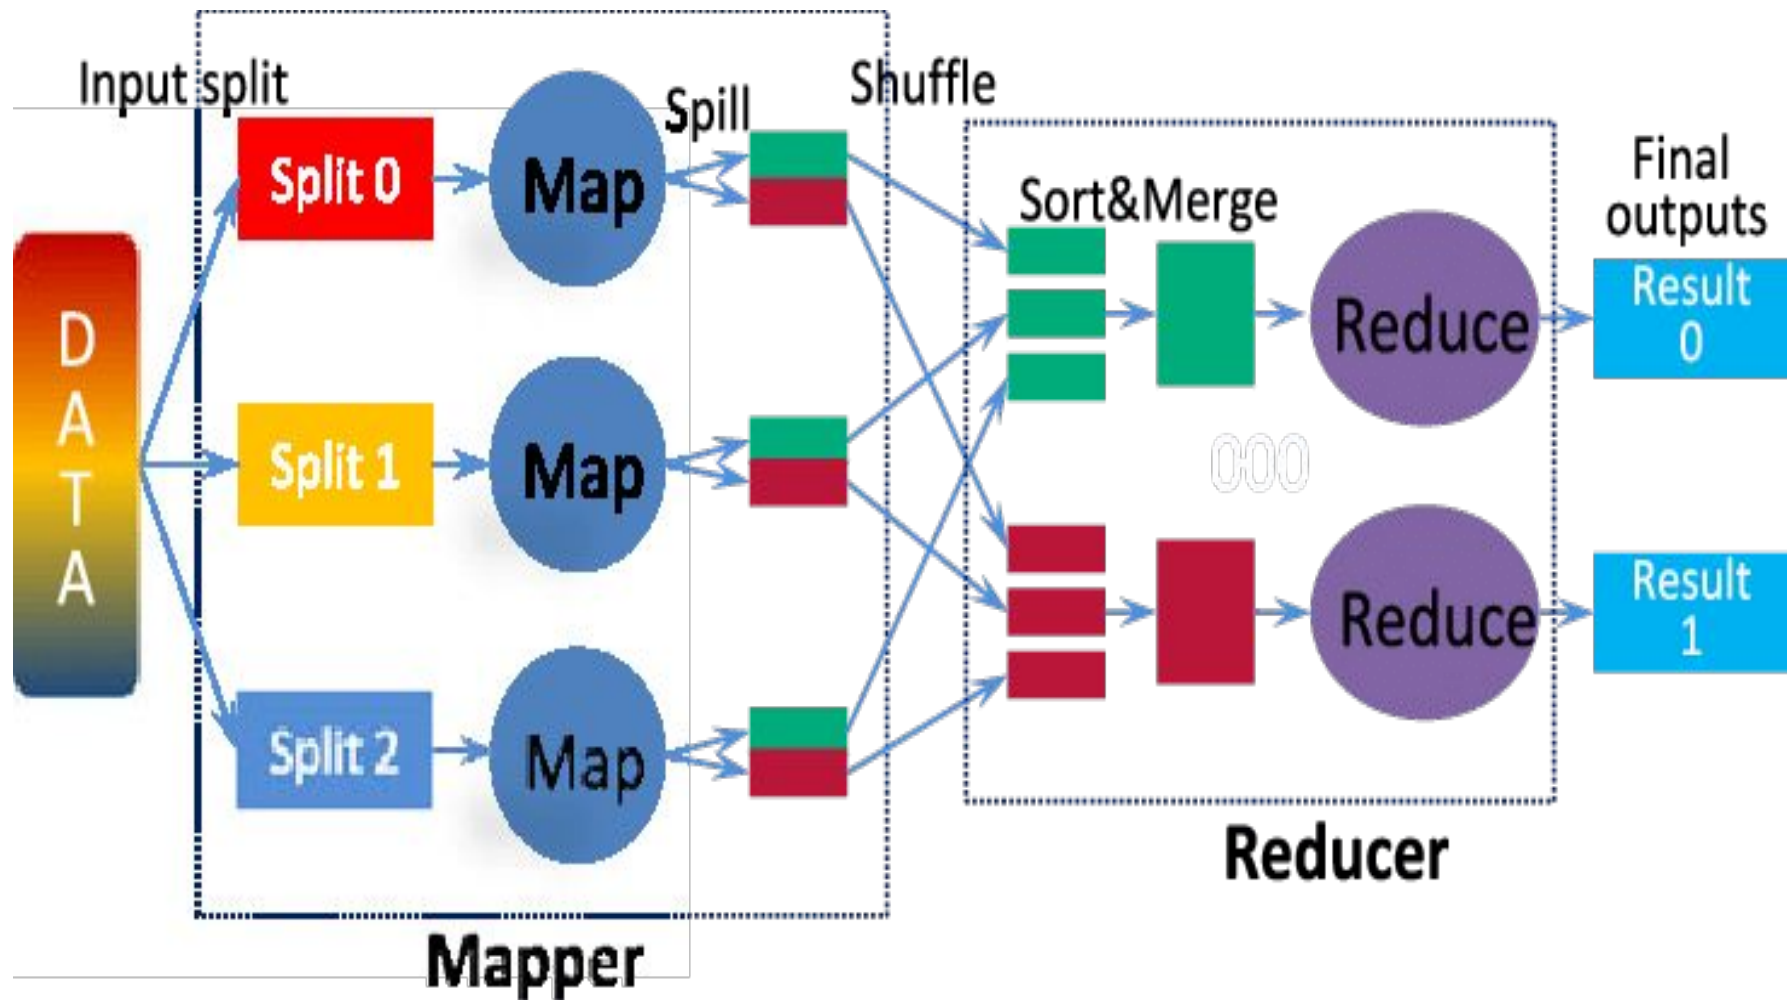
Combiner

- Combiner is an optional function but provides high performance in terms of network bandwidth and disk space.
- It takes intermediate key-value pair provided by mapper and applies user-specific aggregate function to only that mapper.
- It is also known as local reducer.

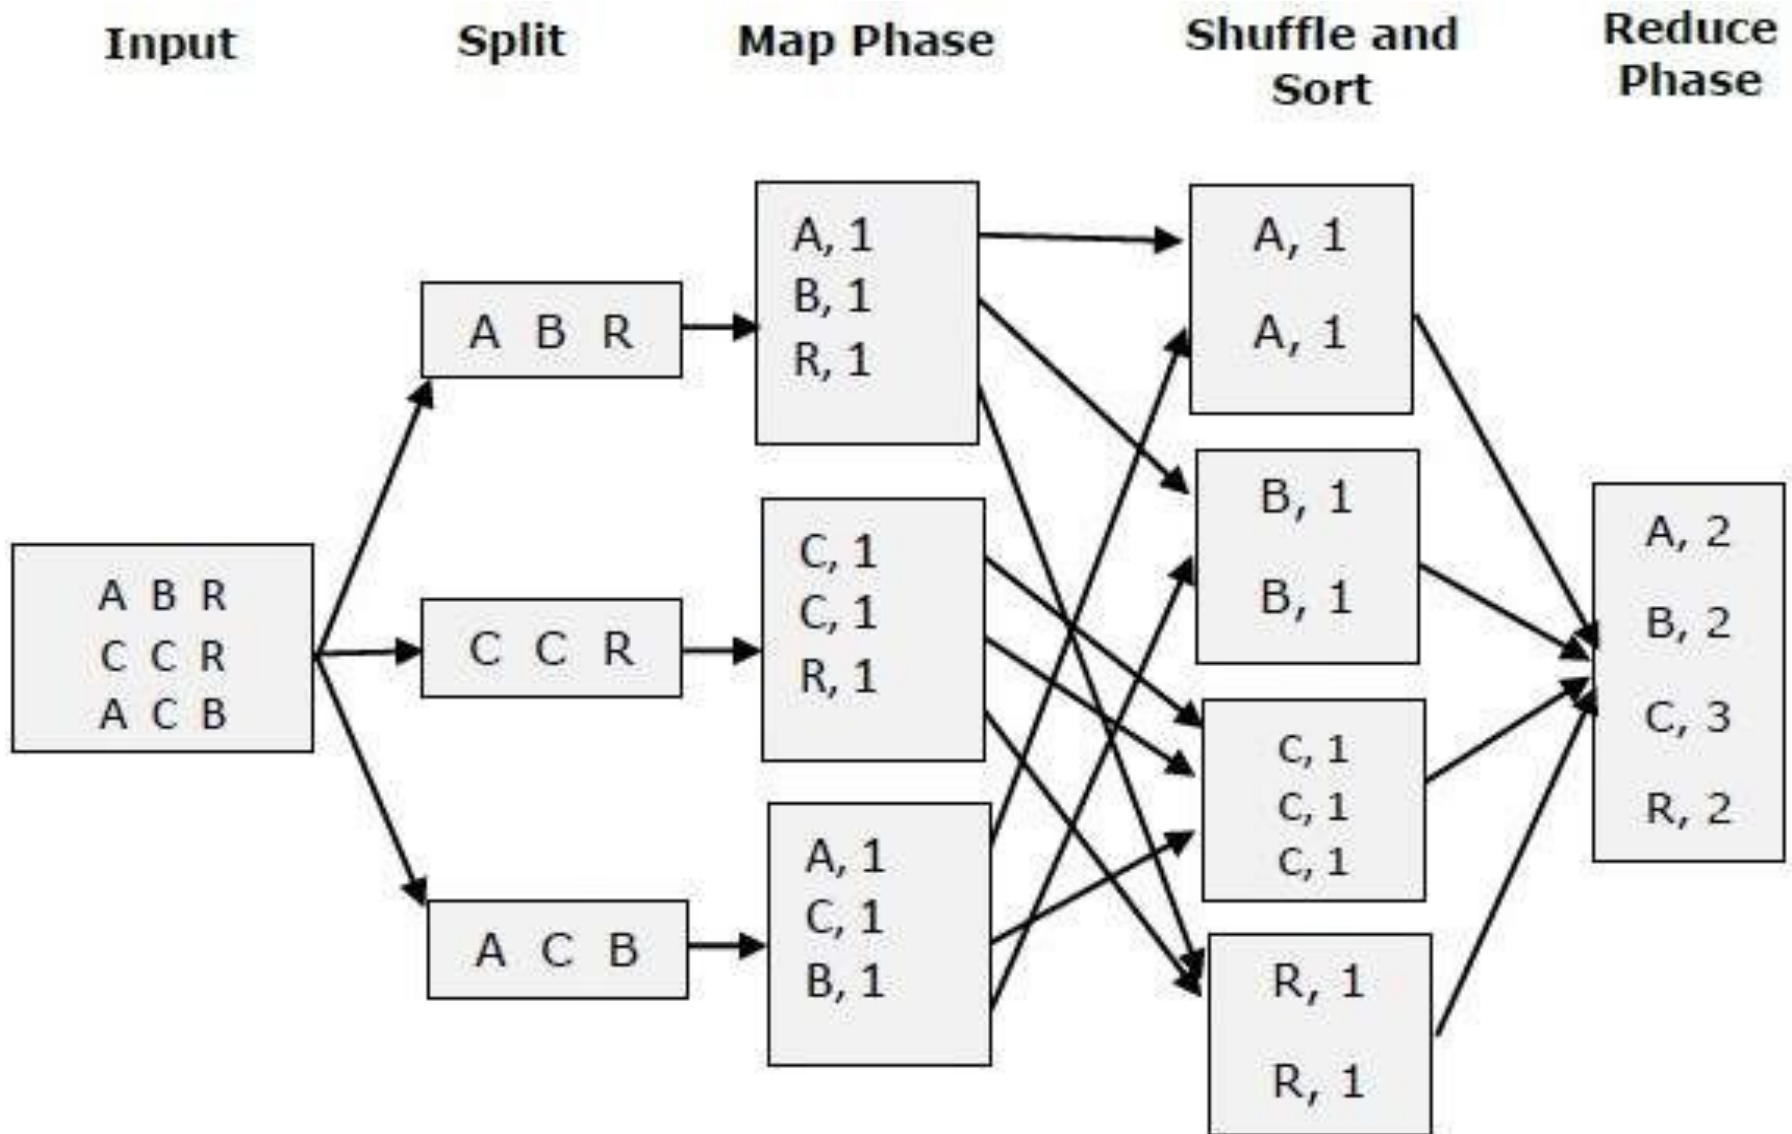
Combiner

- A Combiner, also known as a semi-reducer, is an **optional class that operates by accepting the inputs from the Map class and thereafter passing the output key-value pairs to the Reducer class.**
- The main function of a Combiner is to summarize the map output records with the same key.
- The combiner in MapReduce is also known as 'Mini-reducer'.
- The primary job of Combiner is **to process the output data from the Mapper, before passing it to Reducer.**
- It runs after the mapper and before the Reducer and its use is optional.

Overall Diagram

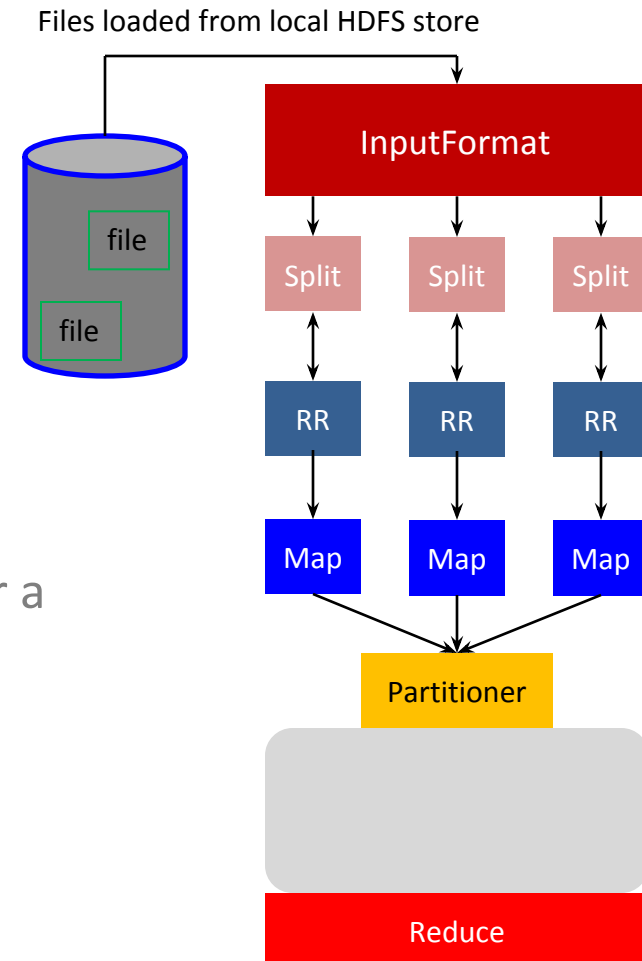


Combiner Task



Partitioner

- Each mapper may emit (K, V) pairs to *any* partition
- Therefore, the map nodes must all agree on where to send different pieces of intermediate data
- The *partitioner* class determines which partition a given (K,V) pair will go to
- The default partitioner computes *a hash value* for a given key and assigns it to a partition based on this result



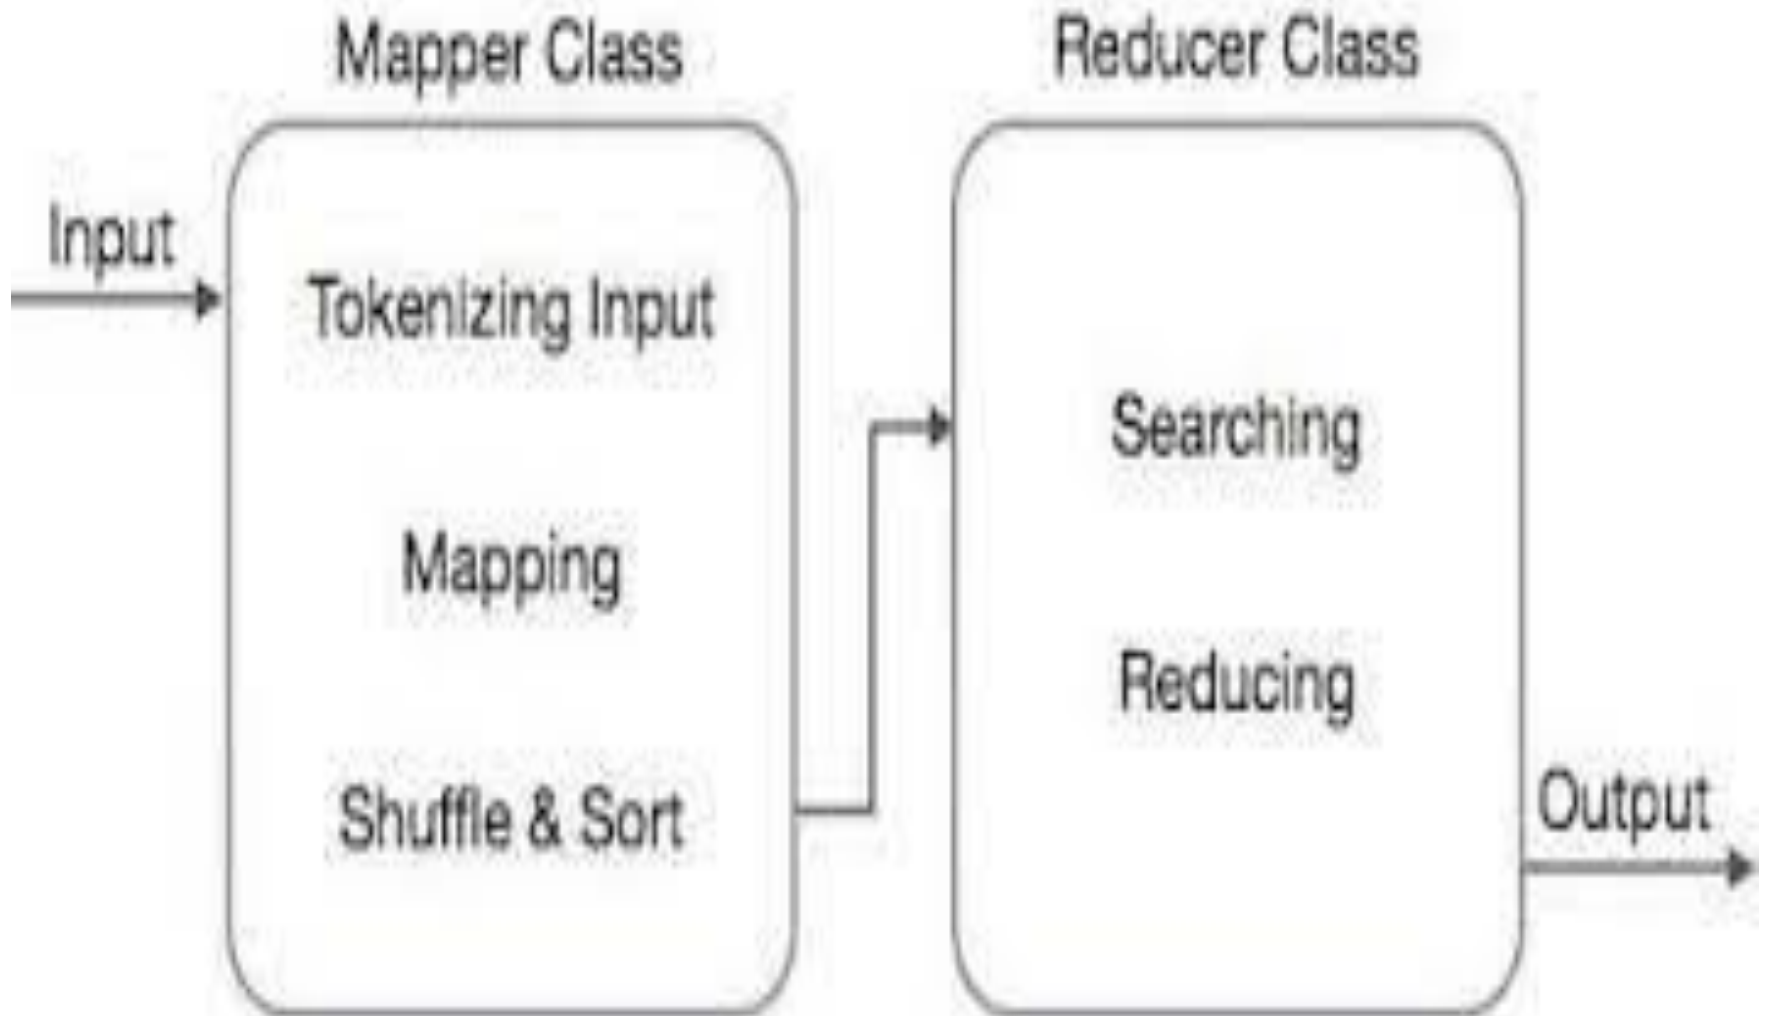
Reducer

- The output produced by map task is known as intermediate keys and values.
- These intermediate keys and values are sent to reducer.
- The reducer tasks are broken into the following phases :
 - 1. Shuffle
 - 2. Sort
 - 3. Reducer
 - 4. output Format

Searching

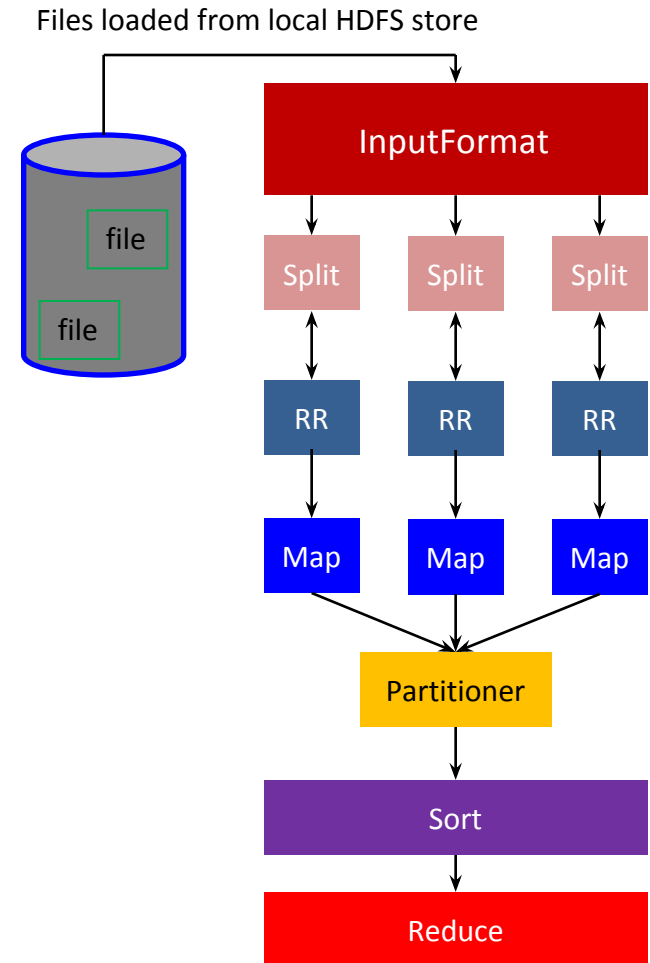
- Searching plays a key role in MapReduce algorithm.
- **It supports in the combiner phase and the Reducer phase.**
- The Map phase processes each input file and provides the data in key-value pairs ($\langle k, v \rangle$).
- The combiner phase (searching technique) will accept the input from the Map phase as a key-value pair.

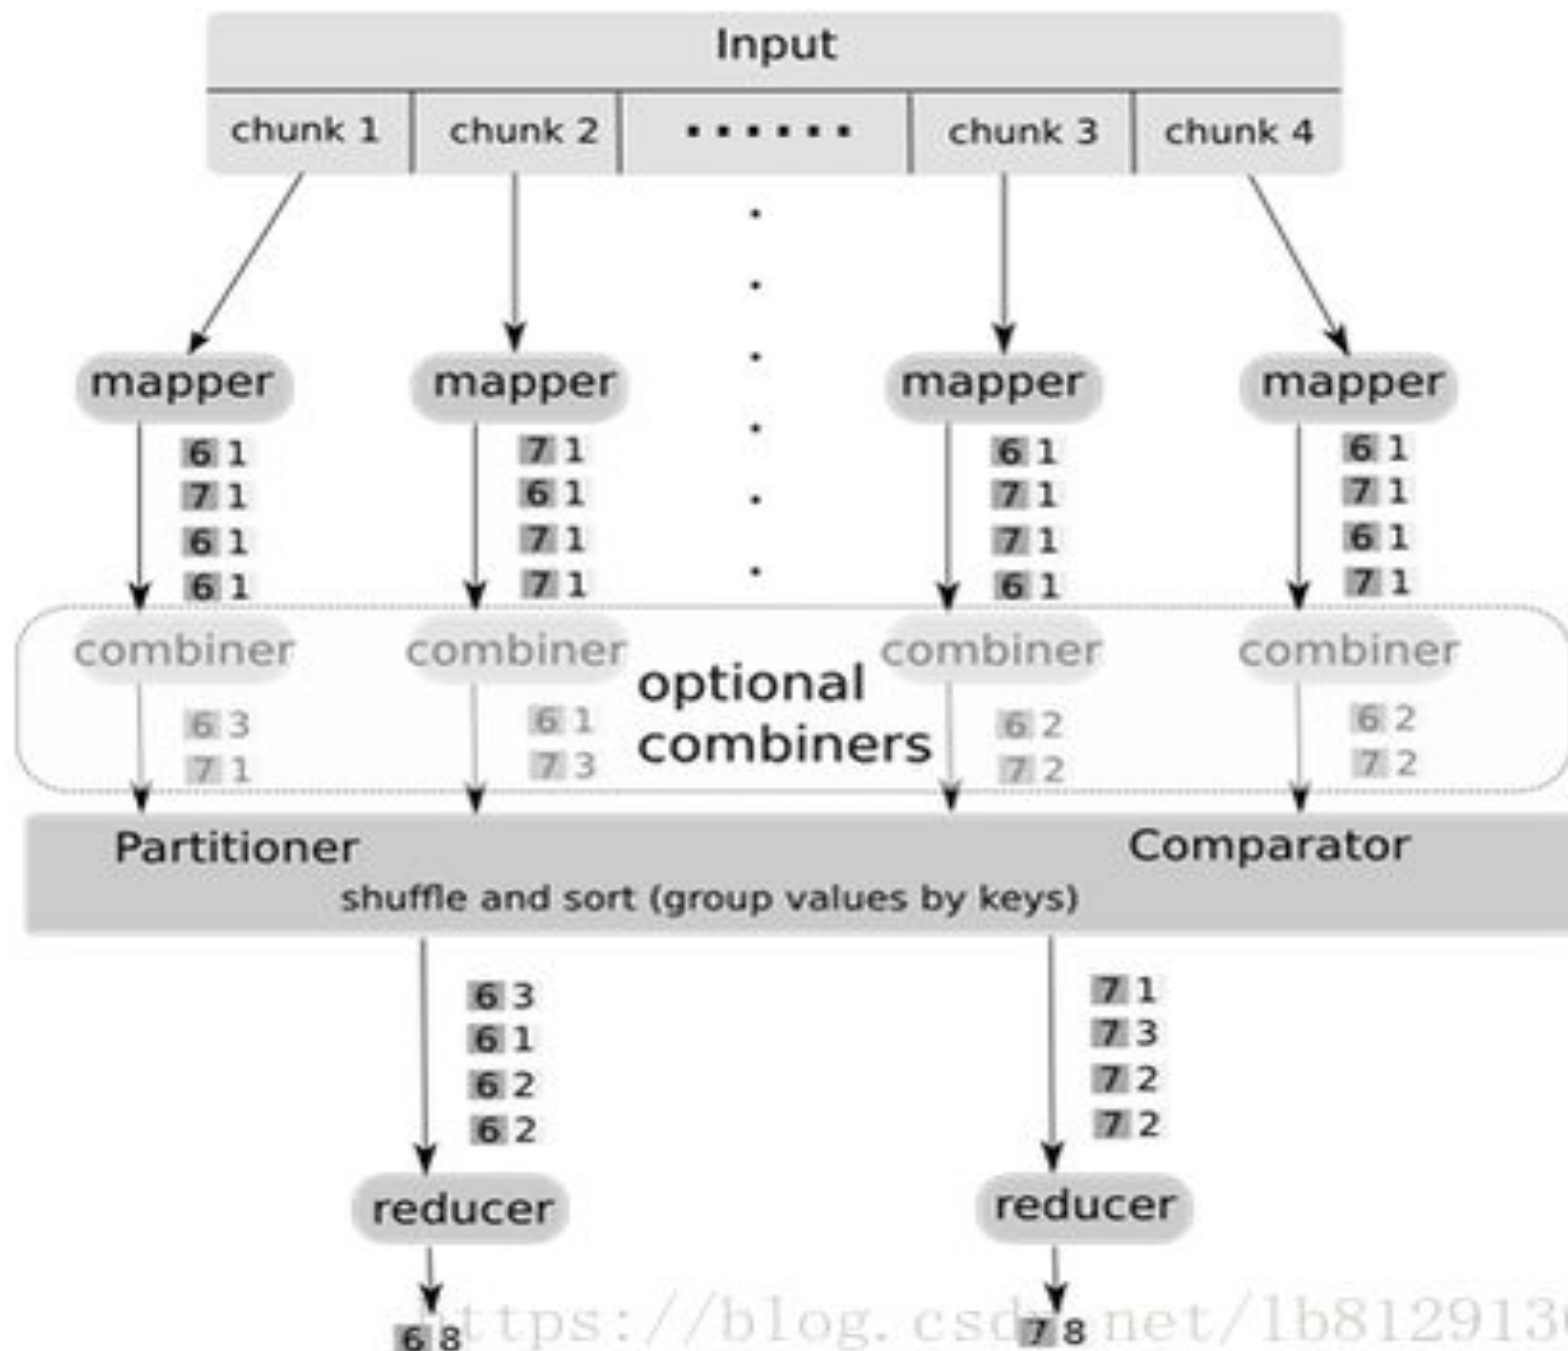
Searching



Sort

- Each Reducer is responsible for reducing the values associated with (several) intermediate keys
- The set of intermediate keys on a single node is *automatically sorted* by MapReduce before they are presented to the Reducer





Compression

- Compress the Mapreduce output files.
- Compression provide two benefits as,
- Reduce the space to store files.
- Speeds up data transfer across the network.
- **Input files are compressed, they will be automatically decompressed.**
- GzipCodec is the compression algorithm for gzip. This compresses the output file.