

# **BIG DATA ANALYTICS LAB**

**(Course Code: 22UPCSC1C11)**

**A programming laboratory record submitted to Periyar University, Salem**

**In partial fulfillment of the requirements for the degree of**

**MASTER OF COMPUTER APPLICATIONS**

**By**

**ELANCHEZHIAN M.**

**[Reg. No.: U22PG507CAP006]**



**DEPARTMENT OF COMPUTER SCIENCE**

**PERIYAR UNIVERSITY**

**(NAAC `A++` Grade with CGPA 3.61) – NIRF RANK 59 – ARIIA RANK 10**

**PERIYAR PALKALAI NAGAR,**

**SALEM – 636 011.**

**(APRIL - 2023)**

## CERTIFICATE

This is to certify that the Programming Laboratory entitled  
“**BIG DATA ANALYTICS - LAB (22UPCSC1C11)**” is a bonafide record  
work done by Mr. /Ms. \_\_\_\_\_  
Register No: \_\_\_\_\_ as partial fulfillment of  
the requirements for the degree of Master of Computer Applications, in the  
Department of Computer Science, Periyar University, Salem, during the  
Academic Year 2023-2024.

Staff In-charge

Head of the Department

Submitted for the practical examination held on.....

Internal Examiner

External Examiner

## **CONTENTS**

<b>S.NO</b>	<b>DATE</b>	<b>TITLE OF THE PROGRAM</b>	<b>PAGE NO</b>	<b>SIGNATURE</b>
1.		Implement File System Shell Commands for HDFS in Hadoop Environment		
2.		Write a Mapreduce program using single reduce function for finding Maximum and Minimum Number		
3.		Write a Mapreduce program using multiple reduce function for Word Count in an given Text document		
4.		Implement the following using Pig Latin Input and Output Operations Relational Operations		
5.		Implement the following using Pig Latin User Defined Functions Advanced Relational Operations		
6.		Write a Word Count program using Pig Latin Script		
7.		Write a program to find a maximum temperature using Pig Latin Script		
8.		Implement the following using Hive commands Handling the Database Creating and Manipulating table		
9.		Implement Simple Queries for database using MongoDB		
10.		Implement Simple Queries for collections using MongoDB		

## **SOURCE CODE:**

**hduser@elan-VirtualBox:/home/elan\$ start-dfs.sh**

Starting namenodes on [localhost]

Starting datanodes

Starting secondary namenodes [elan-VirtualBox]

2023-03-04 22:24:49,572 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

**hduser@elan-VirtualBox:/home/elan\$ start-yarn.sh**

Starting resourcemanager

resourcemanager is running as process 8979. Stop it first and ensure /tmp/hadoop-hduser-resourcemanager.pid file is empty before retry.

Starting nodemanagers

localhost: nodemanager is running as process 9110. Stop it first and ensure /tmp/hadoop-hduser-nodemanager.pid file is empty before retry.

**hduser@elan-VirtualBox:/home/elan\$ jps**

8979 ResourceManager

9110 NodeManager

12888 DataNode

13098 SecondaryNameNode

13535 Jps

12751 NameNode

**hduser@elan-VirtualBox:/home/elan\$ hdfs dfs -ls /**

2023-03-04 22:32:15,190 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Found 1 items

drwx----- - hduser supergroup 0 2023-03-04 22:17 /user

**hduser@elan-VirtualBox:/home/elan\$ hdfs dfs -mkdir /elan**

2023-03-04 22:33:30,869 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

**hduser@elan-VirtualBox:/home/elan\$ hdfs dfs -ls /**

2023-03-04 22:33:37,423 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Found 2 items

drwxr-xr-x	- hduser supergroup	0	2023-03-04 22:33	/elan
drwx-----	- hduser supergroup	0	2023-03-04 22:17	/user

**hduser@elan-VirtualBox:/home/elan\$ hdfs dfs -touchz /elan/magu.txt**

2023-03-04 22:34:15,934 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

**hduser@elan-VirtualBox:/home/elan\$ hdfs dfs -ls /elan/**

2023-03-04 22:34:26,146 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Found 1 items

-rw-r--r--	1 hduser supergroup	0	2023-03-04 22:34	/elan/magu.txt
------------	---------------------	---	------------------	----------------

**hduser@elan-VirtualBox:/home/elan\$ hdfs dfs -put AI.txt /elan/**

2023-03-04 22:36:16,386 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

**hduser@elan-VirtualBox:/home/elan\$ hdfs dfs -ls /elan/**

2023-03-04 22:36:25,867 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Found 2 items

-rw-r--r--	1 hduser supergroup	7	2023-03-04 22:36	/elan/AI.txt
-rw-r--r--	1 hduser supergroup	0	2023-03-04 22:34	/elan/magu.txt

**hduser@elan-VirtualBox:/home/elan\$ hdfs dfs -cat /elan/AI.txt**

2023-03-04 22:38:21,461 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Hellow

**hduser@elan-VirtualBox:/home/elan\$ hdfs dfs -du /elan/AI.txt**

2023-03-04 22:39:50,335 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

7 7 /elan/AI.txt

**hduser@elan-VirtualBox:/home/elan\$ hdfs dfs -dus /user**

2023-03-04 22:40:51,214 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

dus: DEPRECATED: Please use 'du -s' instead.

0 0 /user

**hduser@elan-VirtualBox:/home/elan\$ hdfs dfs -stat /elan/**

2023-03-04 22:41:17,430 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

2023-03-04 17:06:18

**hduser@elan-VirtualBox:/home/elan\$ hdfs dfs -setrep -R 4 /elan**

2023-03-04 22:42:40,864 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Replication 4 set: /elan/AI.txt

Replication 4 set: /elan/magu.txt

**hduser@elan-VirtualBox:/home/elan\$ echo \$?**

0

**hduser@elan-VirtualBox:/home/elan\$ hdfs dfs -test /elan/magu.txt&echo \$?**

[1] 15182

0

**hduser@elan-VirtualBox:/home/elan\$ hdfs dfs -count -q /elan/AI.txt**

2023-03-04 22:46:43,775 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

none	inf	none	inf	0	1	7 /elan/AI.txt
------	-----	------	-----	---	---	----------------

**hduser@elan-VirtualBox:/home/elan\$ hdfs dfs -mkdir /temp**

2023-03-04 22:49:01,972 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

**hduser@elan-VirtualBox:/home/elan\$ hdfs dfs -cp /elan /temp**

2023-03-04 22:49:35,176 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

**hduser@elan-VirtualBox:/home/elan\$ hdfs dfs -ls /temp**

2023-03-04 22:49:50,868 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Found 1 items

drwxr-xr-x	- hduser supergroup	0	2023-03-04 22:49	/temp/elan
------------	---------------------	---	------------------	------------

**hduser@elan-VirtualBox:/home/elan\$ hdfs dfs -ls /**

2023-03-04 22:50:08,561 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Found 3 items

drwxr-xr-x	- hduser supergroup	0	2023-03-04 22:36	/elan
drwxr-xr-x	- hduser supergroup	0	2023-03-04 22:49	/temp
drwx-----	- hduser supergroup	0	2023-03-04 22:17	/user

**hduser@elan-VirtualBox:/home/elan\$ hdfs dfs -mv /elan/AI.txt /temp/AI.txt**

2023-03-04 22:53:23,022 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

**hduser@elan-VirtualBox:/home/elan\$ hdfs dfs -ls /temp/**

2023-03-04 22:53:40,368 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Found 2 items

-rw-r--r--	4	hduser	supergroup	7	2023-03-04 22:36	/temp/ <b>AI.txt</b>
drwxr-xr-x	-	hduser	supergroup	0	2023-03-04 22:49	/temp/ <b>elan</b>

**hduser@elan-VirtualBox:/home/elan\$ hdfs dfs -rmr /temp**

2023-03-04 22:54:30,679 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

rmr: DEPRECATED: Please use '-rm -r' instead.

2023-03-04 22:54:32,095 INFO fs.TrashPolicyDefault: Moved: 'hdfs://localhost:54310/temp' to trash at:

hdfs://localhost:54310/user/hduser/.Trash/Current/temp

**hduser@elan-VirtualBox:/home/elan\$ hdfs dfs -ls /**

2023-03-04 22:54:50,505 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Found 2 items

drwxr-xr-x	-	hduser	supergroup	0	2023-03-04 22:53	/ <b>elan</b>
drwx-----	-	hduser	supergroup	0	2023-03-04 22:17	/ <b>user</b>

**hduser@elan-VirtualBox:/home/elan\$ stop-all.sh**

WARNING: Stopping all Apache Hadoop daemons as hduser in 10 seconds.

WARNING: Use CTRL-C to abort.

Stopping namenodes on [localhost]

Stopping datanodes

Stopping secondary namenodes [elan-VirtualBox]

2023-03-04 23:33:57,243 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Stopping nodemanagers

Stopping resourcemanager



## **SOURCE CODE:**

**hadoop@ubuntu:~\$ start-all.sh**

WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.

WARNING: This is not a recommended production deployment configuration.

WARNING: Use CTRL-C to abort.

Starting namenodes on [0.0.0.0]

Starting datanodes

Starting secondary namenodes [ubuntu]

2023-03-04 21:29:38,334 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Starting resourcemanager

Starting nodemanagers

**hadoop@ubuntu:~\$ mkdir Bda**

**hadoop@ubuntu:~\$ cd Bda**

**hadoop@ubuntu:~/Bda\$ nano minMaxMapper.py**

```
import sys
for line in sys.stdin:
    val = line.strip()
    year = val[0:4]
    temp = val[5:9]
    if temp != '9999':
        print('%s\t%s' % (year, temp))
```

**hadoop@ubuntu:~/Bda\$ nano minMaxReducer.py**

```
from operator import itemgetter
import sys

(last_key, max_val, min_val) = (None, 0, 0)

for line in sys.stdin:
    (key, val) = line.strip().split("\t")

    if last_key and last_key != key:
        print('%s\t%s\t%s' % (last_key, max_val, min_val))
        (last_key, max_val, min_val) = (key, float(val), float(val))
```

```
else:
    (last_key, max_val, min_val) = (key, max(max_val,
float(val)), min(min_val, float(val)))
    if min_val == 0.0:
        min_val = max_val
if last_key:
    print('%s\t%s\t%s' % (last_key, max_val, min_val))
```

**hadoop@ubuntu:~/Bda\$ nano temp.txt**

```
2011  41
2011  72
2012  55
2012  43
2013  51
2013  62
2014  23
2015  45
2016  47
2016  65
2017  43
2017  23
2018  60
2019  65
2022  78
2022  45
2023  22
```

**hadoop@ubuntu:~/Bda\$ ls**

minMaxMapper.py minMaxReducer.py temp.txt

**hadoop@ubuntu:~/Bda\$ cat temp.txt | python minMaxMapper.py |  
python minMaxReducer.py**

2011	72.0	41.0
2012	55.0	43.0
2013	62.0	51.0
2014	23.0	23.0
2015	45.0	45.0
2016	65.0	47.0

2017	43.0	23.0
2018	60.0	60.0
2019	65.0	65.0
2022	78.0	45.0
2023	22.0	22.0

**hadoop@ubuntu:~/Bda\$ hdfs dfs -mkdir /BigData**

2023-03-04 21:44:12,439 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

**hadoop@ubuntu:~/Bda\$ hdfs dfs -ls /**

2023-03-04 21:44:25,881 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Found 2 items

drwxr-xr-x	- hadoop supergroup	0	2023-03-04 21:44	<b>/BigData</b>
drwxr-xr-x	- hadoop supergroup	0	2023-03-04 21:43	/elan

**hadoop@ubuntu:~/Bda\$ hdfs dfs -put minMaxReducer.py /BigData**

2023-03-04 21:45:16,642 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

**hadoop@ubuntu:~/Bda\$ hdfs dfs -put minMaxMapper.py /BigData**

2023-03-04 21:45:40,287 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

**hadoop@ubuntu:~/Bda\$ hdfs dfs -put temp.txt /BigData**

2023-03-04 21:46:10,971 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

**hadoop@ubuntu:~/Bda\$ hdfs dfs -ls /BigData/**

2023-03-04 21:46:26,374 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Found 3 items

```
-rw-r--r--  1 hadoop supergroup      161 2023-03-04 21:45
/BigData/minMaxMapper.py
-rw-r--r--  1 hadoop supergroup      554 2023-03-04 21:45
/BigData/minMaxReducer.py
-rw-r--r--  1 hadoop supergroup      136 2023-03-04 21:46
/BigData/temp.txt
```

```
hadoop@ubuntu:~/Bda$ hadoop jar /usr/local/hadoop/share/hadoop/
tools/lib/hadoop-streaming-*.jar -file minMaxMapper.py -mapper
"python minMaxMapper.py" -file minMaxReducer.py -reducer "python
minMaxReducer.py" -input /BigData/temp.txt -output /BigData/
P02RESULT
```

```
hadoop@ubuntu:~/Bda$ hdfs dfs -ls /BigData/
```

```
2023-03-04 22:01:01,336 WARN util.NativeCodeLoader: Unable to load
native-hadoop library for your platform... using builtin-java classes where
applicable
```

Found 4 items

```
drwxr-xr-x  - hadoop supergroup        0 2023-03-04 21:58
/BigData/P02RESULT
-rw-r--r--  1 hadoop supergroup      161 2023-03-04 21:45
/BigData/minMaxMapper.py
-rw-r--r--  1 hadoop supergroup      554 2023-03-04 21:45
/BigData/minMaxReducer.py
-rw-r--r--  1 hadoop supergroup      136 2023-03-04 21:46
/BigData/temp.txt
```

```
hadoop@ubuntu:~/Bda$ hdfs dfs -ls /BigData/P02RESULT/
```

```
2023-03-04 22:02:41,070 WARN util.NativeCodeLoader: Unable to load
native-hadoop library for your platform... using builtin-java classes where
applicable
```

Found 2 items

```
-rw-r--r--  1 hadoop supergroup        0 2023-03-04 21:58
/BigData/P02RESULT/_SUCCESS
-rw-r--r--  1 hadoop supergroup      165 2023-03-04 21:58
/BigData/P02RESULT/part-00000
```

**hadoop@ubuntu:~/Bda\$ hdfs dfs -cat /BigData/P02RESULT/part-\***

2023-03-04 22:03:05,017 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

2011	72.0	41.0
2012	55.0	43.0
2013	62.0	51.0
2014	23.0	23.0
2015	45.0	45.0
2016	65.0	47.0
2017	43.0	23.0
2018	60.0	60.0
2019	65.0	65.0
2022	78.0	45.0
2023	22.0	22.0

**hadoop@ubuntu:~/Bda\$ stop-all.sh**

WARNING: Stopping all Apache Hadoop daemons as hadoop in 10 seconds.

WARNING: Use CTRL-C to abort.

Stopping namenodes on [0.0.0.0]

Stopping datanodes

Stopping secondary namenodes [ubuntu]

2023-03-04 22:10:56,169 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Stopping nodemanagers

localhost: WARNING: nodemanager did not stop gracefully after 5 seconds:

Trying to kill with kill -9

Stopping resourcemanager

## SOURCE CODE:

**hadoop@ubuntu:~\$ start-all.sh**

WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.

WARNING: This is not a recommended production deployment configuration.

WARNING: Use CTRL-C to abort.

Starting namenodes on [0.0.0.0]

Starting datanodes

Starting secondary namenodes [ubuntu]

2023-03-05 04:56:47,216 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Starting resourcemanager

Starting nodemanagers

**hadoop@ubuntu:~\$ nano words.txt**

Big data analytics is the use of advanced analytic techniques against very large, diverse data sets that include structured, semi-structured and unstructured data, from different sources, and in different sizes from terabytes to zettabytes

**hadoop@ubuntu:~\$ nano wordMapper.py**

```
import sys
for line in sys.stdin:
    line=line.strip()
    words=line.split()
    for word in words:
        print('%s\t%s'%(word,"1"))
```

**hadoop@ubuntu:~\$ nano wordReducer.py**

```
import sys

wordcount = {}

for line in sys.stdin:
    line = line.strip()
    word, count = line.split()
```

```
try:
    wordcount[word] = wordcount[word] + int(count)
except:
    wordcount[word] = int(count)
```

```
for word in wordcount.keys():
    print('%s\t%s' % (word, wordcount[word]))
```

**hadoop@ubuntu:~\$ hdfs dfs -mkdir -p /Pro3**

2023-03-05 05:19:47,272 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

**hadoop@ubuntu:~\$ hdfs dfs -ls /**

2023-03-05 05:20:00,556 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Found 4 items

drwxr-xr-x	- hadoop supergroup	0	2023-03-04 21:57	/BigData
drwxr-xr-x	- hadoop supergroup	0	2023-03-05 05:19	<b>/Pro3</b>
drwxr-xr-x	- hadoop supergroup	0	2023-03-04 21:43	/elan
drwx-----	- hadoop supergroup	0	2023-03-04 21:52	/tmp

**hadoop@ubuntu:~\$ hdfs dfs -chmod 700 /Pro3**

2023-03-05 05:20:26,829 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

**hadoop@ubuntu:~\$ hdfs dfs -put words.txt /Pro3**

2023-03-05 05:21:14,052 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

**hadoop@ubuntu:~\$ hdfs dfs -put wordMapper.py /Pro3**

2023-03-05 05:21:33,014 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

**hadoop@ubuntu:~\$ hdfs dfs -put wordReducer.py /Pro3**

2023-03-05 05:21:47,709 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

**hadoop@ubuntu:~\$ hdfs dfs -ls /Pro3**

2023-03-05 05:22:00,937 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Found 3 items

```
-rw-r--r--  1 hadoop supergroup      159 2023-03-05 05:21
/Pro3/wordMapper.py
-rw-r--r--  1 hadoop supergroup      342 2023-03-05 05:21
/Pro3/wordReducer.py
-rw-r--r--  1 hadoop supergroup      242 2023-03-05 05:21 /Pro3/words.txt
```

**hadoop@ubuntu: ~\$ cat words.txt | python wordMapper.py | sort |  
python wordReducer.py**

and	2
structured,	1
semi-structured	1
is	1
terabytes	1
analytics	1
diverse	1
in	1
techniques	1
data,	1
different	2
from	2
to	1
Big	1
analytic	1
large,	1
include	1
that	1
very	1



use	1
unstructured	1
data	2
advanced	1
zettabytes	1
sources,	1
sizes	1
of	1
against	1
sets	1
the	1

```
hadoop@ubuntu: ~$ hadoop jar /usr/local/hadoop/share/hadoop/
tools/lib/hadoop-streaming-*.jar -D mapred.reduce.tasks=2 -file
wordMapper.py -mapper "python wordMapper.py" -file
wordReducer.py -reducer "python wordReducer.py" -input /Pro3/
words.txt -output /Pro3/WordCoutResult
```

```
hadoop@ubuntu: ~$ hdfs dfs -ls /Pro3/WordCoutResult
```

```
2023-03-05 05:30:11,919 WARN util.NativeCodeLoader: Unable to load
native-hadoop library for your platform... using builtin-java classes where
applicable
```

```
Found 3 items
```

```
-rw-r--r--  1 hadoop supergroup      0 2023-03-05 05:29
/Pro3/WordCoutResult/_SUCCESS
-rw-r--r--  1 hadoop supergroup    141 2023-03-05 05:29
/Pro3/WordCoutResult/part-00000
-rw-r--r--  1 hadoop supergroup    135 2023-03-05 05:29
/Pro3/WordCoutResult/part-00001
```

```
hadoop@ubuntu:~$ hdfs dfs -cat /Pro3/WordCoutResult/part*
```

```
2023-03-05 05:31:28,786 WARN util.NativeCodeLoader: Unable to load
native-hadoop library for your platform... using builtin-java classes where
applicable
```

and	2
different	2
zettabytes	1
structured,	1

to	1
that	1
of	1
terabytes	1
sets	1
use	1
against	1
analytic	1
large,	1
in	1
the	1
techniques	1
data,	1
from	2
sources,	1
sizes	1
very	1
Big	1
is	1
unstructured	1
analytics	1
diverse	1
semi-structured	1
include	1
data	2
advanced	1

**hadoop@ubuntu:~\$ stop-all.sh**

WARNING: Stopping all Apache Hadoop daemons as hadoop in 10 seconds.

WARNING: Use CTRL-C to abort.

Stopping namenodes on [0.0.0.0]

Stopping datanodes

Stopping secondary namenodes [ubuntu]

2023-03-05 05:32:11,114 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Stopping nodemanagers

Stopping resourcemanager

## **SOURCE CODE:**

```
test> use BDA
```

```
BDA> db.stuInfo.insertOne(  
  {"_id" : 1, Name:'Elanchezhian', Age:'21', Course:'MCA'}  
)
```

```
BDA> db.stuInfo.insertMany([  
  {"_id" : 2, Name : 'Dhanush', Age : '22', Course:'MSC CS'},  
  {"_id" : 3, Name : 'Hariharan', Age : '20', Course: 'MSC DS'},  
  {"_id" : 4, Name : 'Lachu', Age : '24', Course: 'MSc IT'}  
)
```

```
BDA> db.stuInfo.find()
```

```
BDA> db.stuInfo.updateOne({ "_id":3 }, { $set: { Course: 'MSc IT' } })
```

```
BDA> db.stuInfo.updateMany({},{$set:{Batch:'2022'}})
```

```
BDA> db.stuInfo.find()
```

```
BDA> db.stuInfo.deleteOne({Name:'Lachu'})
```

```
BDA> db.stuinfo.deleteMany({Batch:'2022'})
```

```
BDA> db.stuInfo.find()
```

## OUTPUT:

- Insert Single Documents:

```
{ acknowledged: true, insertedId: 1 }
```

- Insert Many Documents:

```
{ acknowledged: true, insertedIds: { '0': 2, '1': 3, '2': 4 } }
```

- Database:

```
[  
  { _id: 1, Name: 'Elanchezhian', Age: '21', Course: 'MCA' },  
  { _id: 2, Name: 'Dhanush', Age: '22', Course: 'MSC CS' },  
  { _id: 3, Name: 'Hariharan', Age: '20', Course: 'MSC DS' },  
  { _id: 4, Name: 'Lachu', Age: '24', Course: 'MSc IT' }  
]
```

- Update Single Document:

```
{  
  acknowledged: true,  
  insertedId: null,  
  matchedCount: 1,  
  modifiedCount: 1,  
  upsertedCount: 0  
}
```

- Update Multiple Document:

```
{  
  acknowledged: true,  
  insertedId: null,  
  matchedCount: 4,  
  modifiedCount: 4,  
  upsertedCount: 0  
}
```

➤ Database:

```
[
  {
    _id: 1,
    Name: 'Elanchezhian',
    Age: '21',
    Course: 'MCA',
    Batch: '2022'
  },
  {
    _id: 2,
    Name: 'Dhanush',
    Age: '22',
    Course: 'MSC CS',
    Batch: '2022'
  },
  {
    _id: 3,
    Name: 'Hariharan',
    Age: '20',
    Course: 'MSc IT',
    Batch: '2022'
  },
  { _id: 4, Name: 'Lachu', Age: '24', Course: 'MSc IT', Batch: '2022' }
]
```

➤ Delete Single Documents:

```
{ acknowledged: true, deletedCount: 1 }
```

➤ Delete Many Documents:

```
{ acknowledged: true, deletedCount: 3 }
```

## **SOURCE CODE:**

```
BDA > db.stuinfo.insertMany([
  { "_id": 1, Name: 'Divya', Age: '22', Course: 'MSC CS' },
  { "_id": 2, Name: 'Shalini', Age: '20', Course: 'MSC DS' },
  { "_id": 3, Name: 'Hariharan', Age: '24', Course: 'MSc IT' }
])
```

```
BDA > db.stuinfo.find()
```

```
BDA > db.stuinfo.find({Name:'Shalini'})
```

```
BDA > db.stuinfo.find({$or:[{Name:'Divya'},{Age:'20'}]}))
```

```
BDA > db.stuinfo.find({Name:{ $in:["Hariharan","Divya"]}}))
```

```
BDA > db.stuinfo.find({ Age: { $gt:'21'}})
```

```
BDA > db.stuinfo.find({ Age: { $lt:'21'}})
```

```
BDA > db.stuinfo.find({ Name: { $eq:'Shalini'}})
```

```
BDA > db.stuinfo.find({ Name: { $nin:['Shalini']}})
```

## OUTPUT:

### ➤ Finding Documents:

```
[
  { _id: 1, Name: 'Divya', Age: '22', Course: 'MSC CS' },
  { _id: 2, Name: 'Shalini', Age: '20', Course: 'MSC DS' },
  { _id: 3, Name: 'Hariharan', Age: '24', Course: 'MSc IT' }
]
```

### ➤ AND Operator:

:

```
[ { _id: 2, Name: 'Shalini', Age: '20', Course: 'MSC DS' } ]
```

### ➤ OR Operator:

```
[
  { _id: 1, Name: 'Divya', Age: '22', Course: 'MSC CS' },
  { _id: 2, Name: 'Shalini', Age: '20', Course: 'MSC DS' }
]
```

### ➤ IN Operator:

```
[
  { _id: 1, Name: 'Divya', Age: '22', Course: 'MSC CS' },
  { _id: 3, Name: 'Hariharan', Age: '24', Course: 'MSc IT' }
]
```

### ➤ Greater than Operator:

```
[
  { _id: 1, Name: 'Divya', Age: '22', Course: 'MSC CS' },
  { _id: 3, Name: 'Hariharan', Age: '24', Course: 'MSc IT' }
]
```

### ➤ Less than Operator:

```
[ { _id: 2, Name: 'Shalini', Age: '20', Course: 'MSC DS' } ]
```

➤ Equal to Operator:

```
[ { _id: 2, Name: 'Shalini', Age: '20', Course: 'MSC DS' } ]
```

➤ NIN Operator:

```
[  
  { _id: 1, Name: 'Divya', Age: '22', Course: 'MSC CS' },  
  { _id: 3, Name: 'Hariharan', Age: '24', Course: 'MSc IT' }  
]
```



## **SOURCE CODE:**

**rr@ubuntu:~\$ start-all.sh**

WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.

WARNING: This is not a recommended production deployment configuration.

WARNING: Use CTRL-C to abort.

Starting namenodes on [0.0.0.0]

Starting datanodes

Starting secondary namenodes [ubuntu]

2023-03-04 21:29:38,334 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Starting resourcemanager

Starting nodemanagers

**rr@ubuntu:~\$ cd BDA**

**rr@ubuntu:~/BDA\$ gedit emp.txt**

```
Elanchezhian,20,100000,CEO
Sivakumar,21,50000,Software_Developer
Magudeshwaren,23,10000,Designer
Dhasush,22,80000,HR
Hariharan,23,4000,Cleaner
Sanjai,44,70000,Recptionist
Lachu,24,10000,Security
```

**rr@ubuntu:~/BDA\$ hive**

**hive> CREATE DATABASE IF NOT EXISTS ELAN;**

OK

Time taken: 0.361 seconds

**hive> SHOW DATABASES;**

OK

default

elan

kera\_company

Time taken: 0.571 seconds, Fetched: 3 row(s)

**hive> ALTER DATABASE ELAN SET DBPROPERTIES('CREATED BY'='ELANCHEZHIAN');**

OK

Time taken: 0.295 seconds

**hive> CREATE TABLE EMP\_CAT(EMP\_NAME STRING, AGE INT, SALARY FLOAT,EMP\_ROLL STRING) ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' STORED AS TEXTFILE;**

OK

Time taken: 0.138 seconds

**hive> LOAD DATA LOCAL INPATH '/home/rr/BDA/emp.txt' INTO TABLE EMP\_CAT;**

Loading data to table default.emp\_cat

Table default.emp\_cat stats: [numFiles=1, totalSize=193]

OK

Time taken: 0.424 seconds

**hive> SELECT \* FROM EMP\_CAT;**

OK

Elanchezhian	20	100000.0	CEO
Sivakumar	21	50000.0	Software_Developer
Magudeshwaren	23	10000.0	Designer
Dhasush	22	80000.0	HR
Hariharan	23	4000.0	Cleaner
Sanjai	44	70000.0	Receptionist
Lachu	24	10000.0	Security

Time taken: 0.186 seconds, Fetched: 7 row(s)

**hive> CREATE TABLE EMP\_CAT2 LIKE EMP\_CAT;**

OK

Time taken: 0.177 seconds

**hive> INSERT INTO TABLE EMP\_CAT2 SELECT \* FROM EMP\_CAT;**

Query ID = rr\_20230327001239\_f8973b2d-95eb-4d6d-9b68-473c3c68bd1e

Total jobs = 3

Launching Job 1 out of 3

.

.

OK

Time taken: 27.428 seconds

**hive> SELECT \* FROM EMP\_CAT2;**

OK

Elanchezhian	20	100000.0	CEO
Sivakumar	21	50000.0	Software_Developer
Magudeshwaren	23	10000.0	Designer
Dhasush	22	80000.0	HR
Hariharan	23	4000.0	Cleaner
Sanjai	44	70000.0	Recptionist
Lachu	24	10000.0	Security

Time taken: 0.171 seconds, Fetched: 7 row(s)

**hive> SELECT \* FROM EMP\_CAT2 WHERE AGE > 21;**

OK

Magudeshwaren	23	10000.0	Designer
Dhasush	22	80000.0	HR
Hariharan	23	4000.0	Cleaner
Sanjai	44	70000.0	Recptionist
Lachu	24	10000.0	Security

Time taken: 0.295 seconds, Fetched: 5 row(s)

**hive> SELECT \* FROM EMP\_CAT2 WHERE SALARY >80000;**

OK

Elanchezhian	20	100000.0	CEO
--------------	----	----------	-----

Time taken: 0.159 seconds, Fetched: 1 row(s)

**rr@ubuntu:~\$ stop-all.sh**

WARNING: Stopping all Apache Hadoop daemons as hadoop in 10 seconds.

WARNING: Use CTRL-C to abort.

Stopping namenodes on [0.0.0.0]

Stopping datanodes

Stopping secondary namenodes [ubuntu]

2023-03-05 05:32:11,114 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Stopping nodemanagers

Stopping resourcemanager

## **SOURCE CODE:**

**hadoop@ubuntu:~\$ start-all.sh**

WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.

WARNING: This is not a recommended production deployment configuration.

WARNING: Use CTRL-C to abort.

Starting namenodes on [0.0.0.0]

Starting datanodes

Starting secondary namenodes [ubuntu]

2023-04-05 07:39:45,309 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Starting resourcemanager

Starting nodemanagers

**hadoop@ubuntu:~\$ gedit wCount.txt**

Pig Represents Big Data as data flows.

Pig is a high-level platform or tool which is used to process the large datasets.

It provides a high-level of abstraction for processing over the MapReduce.

It provides a high-level scripting language, known as Pig Latin which is used to develop the data analysis codes.

**hadoop@ubuntu:~\$ hdfs dfs -put wCount.txt /elan**

2023-04-05 07:44:11,347 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

**hadoop@ubuntu:~\$ hdfs dfs -ls /elan**

2023-04-05 07:44:25,000 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Found 1 items

-rw-r--r-- 1 hadoop supergroup 313 2023-04-05 07:44 /elan/**wCount.txt**

**hadoop@ubuntu:~\$ gedit wCount.pig**

```
Input_Data = LOAD '/elan/wCount.txt' AS (line:chararray);
tozkenized_data = FOREACH Input_Data GENERATE
FLATTEN(TOKENIZE(line)) AS word;
grouped_data = GROUP tozkenized_data BY word;
word_count = FOREACH grouped_data GENERATE group AS word,
COUNT(tozkenized_data) AS count;
DUMP ordered_data;
```

**hadoop@ubuntu:~\$ pig wCount.pig**

```
(a,3)
(It,2)
(as,2)
(is,3)
(of,1)
(or,1)
(to,2)
(Big,1)
(Pig,3)
(for,1)
(the,3)
(Data,1)
(data,2)
(over,1)
(tool,1)
(used,2)
(Latin,1)
(known,1)
(large,1)
(which,2)
(codes.,1)
(flows.,1)
(develop,1)
(process,1)
```

(analysis,1)  
(language,1)  
(platform,1)  
(provides,2)  
(datasets.,1)  
(scripting,1)  
(MapReduce.,1)  
(Represents,1)  
(high-level,3)  
(processing,1)  
(abstraction,1)

**hadoop@ubuntu:~\$ stop-all.sh**

WARNING: Stopping all Apache Hadoop daemons as hadoop in 10 seconds.

WARNING: Use CTRL-C to abort.

Stopping namenodes on [0.0.0.0]

Stopping datanodes

Stopping secondary namenodes [ubuntu]

2023-04-05 07:39:45,309 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Stopping nodemanagers

Stopping resourcemanager

## **SOURCE CODE:**

**hadoop@ubuntu:~\$ start-all.sh**

WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.

WARNING: This is not a recommended production deployment configuration.

WARNING: Use CTRL-C to abort.

Starting namenodes on [0.0.0.0]

Starting datanodes

Starting secondary namenodes [ubuntu]

2023-04-05 07:39:45,309 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Starting resourcemanager

Starting nodemanagers

**hadoop@ubuntu:~\$ gedit minMax.txt**

2000 45

2000 15

2001 26

2002 24

2003 21

2003 36

2004 26

**hadoop@ubuntu:~\$ hdfs dfs -put minMax.txt /BigData**

2023-04-05 08:20:55,357 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable



**hadoop@ubuntu:~\$ hdfs dfs -ls /BigData**

2023-04-05 08:28:38,199 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Found 6 items

```
drwxr-xr-x  - hadoop supergroup      0 2023-03-04 21:58
/BigData/P02RESULT
-rw-r--r--  1 hadoop supergroup     65 2023-04-05 08:20
/BigData/minMax.txt
-rw-r--r--  1 hadoop supergroup    161 2023-03-04 21:45
/BigData/minMaxMapper.py
-rw-r--r--  1 hadoop supergroup    554 2023-03-04 21:45
/BigData/minMaxReducer.py
-rw-r--r--  1 hadoop supergroup    136 2023-03-04 21:46
/BigData/temp.txt
-rw-r--r--  1 hadoop supergroup    131 2023-04-05 07:11
/BigData/wordCount.txt
```

**hadoop@ubuntu:~\$ gedit temp.pig**

```
A = LOAD '/BigData/minMax.txt' using PigStorage (' ') AS
(year:int,temp:int);
B = GROUP A BY year;
C = FOREACH B GENERATE group, MAX(A.temp), MIN(A.temp);
DUMP C;
```

**hadoop@ubuntu:~\$ pig temp.pig**

```
(2000,45,15)
(2001,26,26)
(2002,24,24)
(2003,36,21)
(2004,26,26)
(2005,37,37)
```

**hadoop@ubuntu:~\$ stop-all.sh**

WARNING: Stopping all Apache Hadoop daemons as hadoop in 10 seconds.

WARNING: Use CTRL-C to abort.

Stopping namenodes on [0.0.0.0]

Stopping datanodes

Stopping secondary namenodes [ubuntu]

2023-04-05 07:39:45,309 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Stopping nodemanagers

Stopping resourcemanager

## **SOURCE CODE:**

**hadoop@ubuntu:~\$ start-all.sh**

WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.

WARNING: This is not a recommended production deployment configuration.

WARNING: Use CTRL-C to abort.

Starting namenodes on [0.0.0.0]

Starting datanodes

Starting secondary namenodes [ubuntu]

2023-04-12 06:37:01,144 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Starting resourcemanager

Starting nodemanagers

**hadoop@ubuntu:~\$ gedit employee.txt**

10012,Santhosh,25,35000

10021,Elanchezhian,26,45250

12045,Suresh,31,85000

14502,Karthi,23,65000

18545,Selva,21,15000

**hadoop@ubuntu:~\$ gedit udf.py**

@outputSchema('num:long')

def get\_length(data):

    str\_data = ''.join([chr(x) for x in data])

    return len(str\_data)

**hadoop@ubuntu:~\$ hdfs dfs -put employee.txt /BDA**

2023-04-12 06:43:15,455 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

**hadoop@ubuntu:~\$ hdfs dfs -put udf.py /BDA**

2023-04-12 06:43:26,792 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

**hadoop@ubuntu:~\$ hdfs dfs -ls /BDA**

2023-04-12 06:44:01,322 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Found 2 items

-rw-r--r-- 1 hadoop supergroup 117 2023-04-12 06:43  
/BDA/**employee.txt**

-rw-r--r-- 1 hadoop supergroup 120 2023-04-12 06:43 /BDA/**udf.py**

**hadoop@ubuntu:~\$ pig**

**grunt> REGISTER 'udf.py' USING jython as pyudf**

**grunt> A = LOAD '/BDA/employee.txt' USING PigStorage();**

**grunt> B = FOREACH A GENERATE  
\$0,pyudf.get\_length(\$0);**

**grunt> DUMP B;**

(10012,Santhosh,25,35000,23)

(10021,Elanchezhian,26,45250,27)

(12045,Suresh,31,85000,21)

(14502,Karthi,23,65000,21)

(18545,Selva,21,15000,20)

**grunt> quit**

**hadoop@ubuntu:~\$ gedit emp\_Del.txt**

10012,Salem,MCA,Developer

10021,Erode,MCA,CEO

12045,Trichy,MCA,Manager

14502,Rahul,B.E,GM

18545,Rehan,M.E,MD

**hadoop@ubuntu:~\$ hdfs dfs -put emp\_Del.txt /BDA**

**hadoop@ubuntu:~\$ pig**

**grunt> A = LOAD '/BDA/employee.txt' USING PigStorage(',') AS  
(id:int,name:chararray,age:int,Salary:int);**

**grunt> B = LOAD '/BDA/emp\_Del.txt ' USING PigStorage(',') AS  
(id:int,location:chararray,Qualification:chararray,Designation:chararray);**

**grunt> join\_data = JOIN A by id,B by id;**

(10012,Santhosh,25,35000,10012,Salem,MCA,Developer)

(10021,Elanchezhian,26,45250,10021,Erode,MCA,CEO)

(12045,Suresh,31,85000,12045,Trichy,MCA,Manager)

(14502,Karthi,23,65000,14502,Rahul,B.E,GM)

(18545,Selva,21,15000,18545,Rehan,M.E,MD)

**grunt> SPLIT A INTO X IF age==25, Y IF age<=22;**

**grunt> dump X;**

(10012,Santhosh,25,35000)

**grunt> dump Y;**

(18545,Selva,21,15000)

**grunt> grade\_grp = GROUP A BY name;**

```
grunt> grade_avg = FOREACH grade_grp GENERATE A.name,  
A.age, AVG(A.Salary);
```

```
grunt> dump grade_avg;
```

```
((Selva)),{(21)},15000.0)  
((Karthi)),{(23)},65000.0)  
((Suresh)),{(31)},85000.0)  
((Santhosh)),{(25)},35000.0)  
((Elanchezhian)),{(26)},45250.0)
```

```
grunt> cogroup_data = COGROUP A BY id,B BY id;
```

```
grunt> dump cogroup_data;
```

```
(10012,{(10012,Santhosh,25,35000)},{(10012,Salem,MCA,Developer)})  
(10021,{(10021,Elanchezhian,26,45250)},{(10021,Erode,MCA,CEO)})  
(12045,{(12045,Suresh,31,85000)},{(12045,Trichy,MCA,Manager)})  
(14502,{(14502,Karthi,23,65000)},{(14502,Rahul,B.E,GM)})  
(18545,{(18545,Selva,21,15000)},{(18545,Rehan,M.E,MD)})
```

```
grunt> M = FOREACH A GENERATE TOMAP(name, age);
```

```
grunt> dump M;
```

```
([Santhosh#25])  
([Elanchezhian#26])  
([Suresh#31])  
([Karthi#23])  
([Selva#21])
```

```
grunt> quit;
```

**hadoop@ubuntu:~\$ stop-all.sh**

WARNING: Stopping all Apache Hadoop daemons as hadoop in 10 seconds.

WARNING: Use CTRL-C to abort.

Stopping namenodes on [0.0.0.0]

Stopping datanodes

Stopping secondary namenodes [ubuntu]

2023-04-12 07:24:45,390 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Stopping nodemanagers

Stopping resourcemanager

## **SOURCE CODE:**

**hadoop@ubuntu:~\$ start-all.sh**

WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.

WARNING: This is not a recommended production deployment configuration.

WARNING: Use CTRL-C to abort.

Starting namenodes on [0.0.0.0]

Starting datanodes

Starting secondary namenodes [ubuntu]

2023-04-12 07:26:04,750 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Starting resourcemanager

Starting nodemanagers

**hadoop@ubuntu:~\$ gedit dels.txt**

1,Elan,25,A

2,Dhanush,22,B

3,Lachu,30,C

4,Sanjai,27,B

5,Akil,24,A

**hadoop@ubuntu:~\$ hdfs dfs -mkdir /BDA**

2023-04-12 06:42:59,540 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable



**hadoop@ubuntu:~\$ hdfs dfs -chmod 755 /BDA**

2023-04-12 06:43:37,776 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

**hadoop@ubuntu:~\$ hdfs dfs -put dels.txt /BDA**

2023-04-12 07:34:08,843 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

**hadoop@ubuntu:~\$ hdfs dfs -ls /BDA**

2023-04-12 07:34:30,977 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

Found 4 items

-rw-r--r-- 1 hadoop supergroup 66 2023-04-12 07:24 /BDA/dels.txt

**hadoop@ubuntu:~\$ pig**

**grunt> A = LOAD '/BDA/dels.txt' USING PigStorage(',') as  
(id:int,name:chararray,age:int,grade:chararray);**

**grunt> DUMP A;**

(1,Elan,25,A)

(2,Dhanush,22,B)

(3,Lachu,30,C)

(4,Sanjai,27,B)

(5,Akil,24,A)

**grunt> fewrecords = LIMIT A 2;**

**grunt> DUMP fewrecords;**

(1,Elan,25,A)

(2,Dhanush,22,B)

```
grunt> agelimit = FILTER A BY age>25;
```

```
grunt> DUMP agelimit;
```

```
(3,Lachu,30,C)
```

```
(4,Sanjai,27,B)
```

```
grunt> DESCRIBE A;
```

```
A: {id: int,name: chararray,age: int,grade: chararray }
```

```
grunt> groupbyage = GROUP A BY age;
```

```
grunt> DUMP groupbyage;
```

```
(22,{{(2,Dhanush,22,B)}})
```

```
(24,{{(5,Akil,24,A)}})
```

```
(25,{{(1,Elan,25,A)}})
```

```
(27,{{(4,Sanjai,27,B)}})
```

```
(30,{{(3,Lachu,30,C)}})
```

```
grunt> maxage = FOREACH groupbyage GENERATE(A.id,  
A.age), MAX(A.age);
```

```
grunt> DUMP maxage;
```

```
(({(2)},{{(22)}}),22)
```

```
(({(5)},{{(24)}}),24)
```

```
(({(1)},{{(25)}}),25)
```

```
(({(4)},{{(27)}}),27)
```

```
(({(3)},{{(30)}}),30)
```

```
grunt> allagevalues = GROUP A ALL;
```

```
grunt> DUMP allagevalues;
```

```
(all,{{(5,Akil,24,A),(4,Sanjai,27,B),(3,Lachu,30,C),(2,Dhan  
ush,22,B),(1,Elan,25,A)}})
```

```
grunt> sumofage = FOREACH allagevalues GENERATE  
group,SUM(A.age);
```

```
grunt> dump sumofage;
```

```
(all,128)
```

```
grunt> dist = DISTINCT A;
```

```
grunt> DUMP dist;
```

```
(1,Elan,25,A)
```

```
(2,Dhanush,22,B)
```

```
(3,Lachu,30,C)
```

```
(4,Sanjai,27,B)
```

```
(5,Akil,24,A)
```

```
grunt> quit
```

```
hadoop@ubuntu:~$ stop-all.sh
```

```
WARNING: Stopping all Apache Hadoop daemons as hadoop in 10 seconds.
```

```
WARNING: Use CTRL-C to abort.
```

```
Stopping namenodes on [0.0.0.0]
```

```
Stopping datanodes
```

```
Stopping secondary namenodes [ubuntu]
```

```
2023-04-12 08:11:41,223 WARN util.NativeCodeLoader: Unable to load  
native-hadoop library for your platform... using builtin-java classes where  
applicable
```

```
Stopping nodemanagers
```

```
Stopping resourcemanager
```